

Phase 1 — LLM policy discovery (once per project, frozen)

Dataset / repo

Allocator profiler

Strategist

Judge tuner

Grep verification + frozen profile

zero LLM calls on the
measured path

Phase 2 — Deterministic engine (no LLM decision)

candidateScan (frozen
allocator manifest)

Stage A — static fan-out sub-agents (5
content-capable tools)

Stage B — pinned dynamic recipe
(buildTarget → IsanRun)

Stage C — synthesis
(dynamicCoverage stamp)

escalation
only
borderline

Phase 3 — Hybrid judging

Track A — heuristic judge (every bundle)

Track B — LLM judge (borderline only;
consensus opt-in $K > 1$)

Reports and verdicts
(results/)

any failure
→ heuristic
verdict

